

Heart Attack Prediction Using Random Forest Supervised Machine Learning

April 29, 2025

```
[222]: # Importing the necessary libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, confusion_matrix, \
    classification_report
from sklearn.tree import DecisionTreeClassifier

# Loading dataset into Jupyter Notebook
data = pd.read_csv('heart-attack-risk-prediction-dataset.csv')

# Exploring dataset
print(data.describe())

print(data.info())

# There is 9651 patients in the dataset.
# Some of the the patients data is not complete so there is less data to \
    classify their risk levels.
# For this we will replace the null entry with the average in the data set of \
    what is observed
```

	Age	Cholesterol	Heart rate	Diabetes	Family History \
count	9651.000000	9651.000000	9651.000000	9377.000000	9377.000000
mean	0.450254	0.499780	0.050756	0.652554	0.488749
std	0.231154	0.284461	0.024922	0.476184	0.499900
min	0.000000	0.000000	0.000000	0.000000	0.000000
25%	0.258427	0.264286	0.035289	0.000000	0.000000
50%	0.460674	0.499780	0.050412	1.000000	0.000000
75%	0.640449	0.739286	0.065995	1.000000	1.000000
max	1.000000	1.000000	1.000000	1.000000	1.000000

	Smoking	Obesity	Alcohol Consumption	Exercise Hours Per Week \
count	9377.000000	9377.000000	9377.000000	9651.000000
mean	0.902421	0.500160	0.600192	0.502110
std	0.296761	0.500027	0.489885	0.284830
min	0.000000	0.000000	0.000000	0.000000
25%	1.000000	0.000000	0.000000	0.259793
50%	1.000000	1.000000	1.000000	0.502110
75%	1.000000	1.000000	1.000000	0.747086
max	1.000000	1.000000	1.000000	1.000000

	Diet ...	Triglycerides	Physical Activity Days Per Week \
count	9651.000000 ...	9651.000000	9377.000000
mean	1.057093 ...	0.503603	3.501866
std	0.868418 ...	0.286183	2.283833
min	0.000000 ...	0.000000	0.000000
25%	0.000000 ...	0.262338	2.000000
50%	1.000000 ...	0.503603	3.000000
75%	2.000000 ...	0.748052	6.000000
max	3.000000 ...	1.000000	7.000000

	Sleep Hours Per Day	Heart Attack Risk (Binary)	Blood sugar \
count	9651.000000	9651.000000	9651.000000
mean	0.504621	0.345146	0.227018
std	0.327482	0.475440	0.075577
min	0.000000	0.000000	0.000000
25%	0.166667	0.000000	0.227018
50%	0.500000	0.000000	0.227018
75%	0.833333	1.000000	0.227018
max	1.000000	1.000000	1.000000

	CK-MB	Troponin	Heart Attack Risk (Text) \
count	9651.000000	9651.000000	9651.000000
mean	0.048229	0.036512	0.398093
std	0.075959	0.059556	0.737488
min	0.000000	0.000000	0.000000
25%	0.048229	0.036512	0.000000
50%	0.048229	0.036512	0.000000
75%	0.048229	0.036512	0.000000
max	1.000000	1.000000	2.000000

	Systolic blood pressure	Diastolic blood pressure
count	9651.000000	9651.000000
mean	0.449982	0.497553
std	0.170344	0.172033
min	0.000000	0.000000
25%	0.303226	0.348837
50%	0.445161	0.500000
75%	0.600000	0.651163

max 1.000000 1.000000

[8 rows x 26 columns]

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 9651 entries, 0 to 9650

Data columns (total 27 columns):

#	Column	Non-Null Count	Dtype
0	Age	9651 non-null	float64
1	Cholesterol	9651 non-null	float64
2	Heart rate	9651 non-null	float64
3	Diabetes	9377 non-null	float64
4	Family History	9377 non-null	float64
5	Smoking	9377 non-null	float64
6	Obesity	9377 non-null	float64
7	Alcohol Consumption	9377 non-null	float64
8	Exercise Hours Per Week	9651 non-null	float64
9	Diet	9651 non-null	int64
10	Previous Heart Problems	9377 non-null	float64
11	Medication Use	9377 non-null	float64
12	Stress Level	9377 non-null	float64
13	Sedentary Hours Per Day	9651 non-null	float64
14	Income	9651 non-null	float64
15	BMI	9651 non-null	float64
16	Triglycerides	9651 non-null	float64
17	Physical Activity Days Per Week	9377 non-null	float64
18	Sleep Hours Per Day	9651 non-null	float64
19	Heart Attack Risk (Binary)	9651 non-null	float64
20	Blood sugar	9651 non-null	float64
21	CK-MB	9651 non-null	float64
22	Troponin	9651 non-null	float64
23	Heart Attack Risk (Text)	9651 non-null	int64
24	Gender	9651 non-null	object
25	Systolic blood pressure	9651 non-null	float64
26	Diastolic blood pressure	9651 non-null	float64

dtypes: float64(24), int64(2), object(1)

memory usage: 2.0+ MB

None

```
[223]: #Clean dataset by replacing null entries with the mean values.  
data.fillna(data.median(numeric_only=True), inplace=True)
```

```
[224]: Checking for Duplicate rows  
print(data.duplicated().sum())  
  
Ensuring correct data type  
print(data.dtypes)
```

Cell In[224], line 1

Checking for Duplicate rows

~

SyntaxError: invalid syntax

```
[225]: data['Gender'] = data['Gender'].map({'Female': 0, 'Male': 1})
```

```
# See the problematic rows
```

```
print(data[data['Gender'].isna()])
```

```
# Then drop them (recommended)
```

```
data = data.dropna(subset=['Gender'])
```

```
# Check again
```

```
print(data['Gender'].unique())
```

	Age	Cholesterol	Heart rate	Diabetes	Family History	Smoking	\
9377	0.516854	0.49978	0.065995	1.0	0.0	1.0	
9378	0.516854	0.49978	0.065995	1.0	0.0	1.0	
9379	0.516854	0.49978	0.065995	1.0	0.0	1.0	
9380	0.382022	0.49978	0.105408	1.0	0.0	1.0	
9381	0.370787	0.49978	0.042163	1.0	0.0	1.0	
...	
9646	0.404494	0.49978	0.091659	1.0	0.0	1.0	
9647	0.235955	0.49978	0.049496	1.0	0.0	1.0	
9648	0.348315	0.49978	0.088909	1.0	0.0	1.0	
9649	0.370787	0.49978	0.067828	1.0	0.0	1.0	
9650	0.797753	0.49978	0.084326	1.0	0.0	1.0	

	Obesity	Alcohol Consumption	Exercise Hours Per Week	Diet	...	\
9377	1.0		1.0	0.50211	3	...
9378	1.0		1.0	0.50211	3	...
9379	1.0		1.0	0.50211	3	...
9380	1.0		1.0	0.50211	3	...
9381	1.0		1.0	0.50211	3	...
...
9646	1.0		1.0	0.50211	3	...
9647	1.0		1.0	0.50211	3	...
9648	1.0		1.0	0.50211	3	...
9649	1.0		1.0	0.50211	3	...
9650	1.0		1.0	0.50211	3	...

	Physical Activity Days Per Week	Sleep Hours Per Day	\
9377	3.0	0.504621	
9378	3.0	0.504621	

9379	3.0	0.504621
9380	3.0	0.504621
9381	3.0	0.504621
...
9646	3.0	0.504621
9647	3.0	0.504621
9648	3.0	0.504621
9649	3.0	0.504621
9650	3.0	0.504621

	Heart Attack Risk (Binary)	Blood sugar	CK-MB	Troponin \
9377	0.0	0.525692	0.004268	0.000388
9378	0.0	0.112648	0.005002	0.000680
9379	0.0	0.140316	0.024723	0.003690
9380	0.0	0.128458	0.315234	0.000291
9381	0.0	0.482213	1.000000	0.000583
...
9646	0.0	0.081028	0.002029	0.006894
9647	0.0	0.084980	0.003500	0.000194
9648	0.0	0.330040	0.009540	0.000194
9649	0.0	0.197628	0.119858	0.025439
9650	0.0	0.156126	0.006237	0.005923

	Heart Attack Risk (Text)	Gender	Systolic blood pressure \
9377	1	NaN	0.554839
9378	1	NaN	0.458065
9379	2	NaN	0.554839
9380	2	NaN	0.212903
9381	2	NaN	0.445161
...
9646	2	NaN	1.000000
9647	1	NaN	0.445161
9648	1	NaN	0.225806
9649	2	NaN	0.258065
9650	2	NaN	0.322581

	Diastolic blood pressure
9377	0.418605
9378	0.313953
9379	0.418605
9380	0.209302
9381	0.174419
...	...
9646	1.000000
9647	0.186047
9648	0.302326
9649	0.453488
9650	0.313953

```
[274 rows x 27 columns]
[1. 0.]
```

```
[226]: X = data.drop(columns=["Heart Attack Risk (Text)"]) # or whatever your target_
        ↪column is called
y = data["Heart Attack Risk (Text)"]

# Train-test split
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
        ↪random_state=42)

from sklearn.ensemble import RandomForestClassifier

rf_model = RandomForestClassifier(n_estimators=100, random_state=42)
rf_model.fit(X_train, y_train)

if 'Risk_Level' in data.columns:
    data = data.drop(columns=['Risk_Level'])

print(y_train.head())
print(y_train.dtype)
```

```
8093    2
7645    1
5639    0
4689    0
7603    0
Name: Heart Attack Risk (Text), dtype: int64
int64
```

```
[227]: # Making a Random Forest Classifier

from sklearn.metrics import classification_report, accuracy_score

y_pred = rf_model.predict(X_test)
print("Accuracy:", accuracy_score(y_test, y_pred))
print("\nClassification Report:\n", classification_report(y_test, y_pred))
```

```
Accuracy: 0.996268656716418
```

```
Classification Report:
              precision    recall  f1-score   support

0               1.00        1.00        1.00        1476
1               0.96        0.99        0.98         144
```

	2	1.00	0.98	0.99	256
accuracy				1.00	1876
macro avg	0.99	0.99	0.99		1876
weighted avg	1.00	1.00	1.00		1876

[229]: *# Using Bagging Classifier to decrease variance and bias*

```
from sklearn.ensemble import BaggingClassifier
from sklearn.tree import DecisionTreeClassifier

# Base estimator (can be a decision tree or any classifier)
base_estimator = DecisionTreeClassifier(random_state=42)

# Initialize and train Bagging
bagging_model = BaggingClassifier(base_estimator=base_estimator,
    ↪n_estimators=100, random_state=42)
bagging_model.fit(X_train, y_train)

# Predict and evaluate
bagging_preds = bagging_model.predict(X_test)
print("Bagging Classifier Classification Report:")
print(classification_report(y_test, bagging_preds))
print("Accuracy:", accuracy_score(y_test, bagging_preds))
```

/Users/ajastarkey/anaconda3/lib/python3.10/site-packages/sklearn/ensemble/_base.py:166: FutureWarning: `base_estimator` was renamed to `estimator` in version 1.2 and will be removed in 1.4.

warnings.warn(

Bagging Classifier Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	1476
1	0.99	1.00	0.99	144
2	1.00	0.99	1.00	256
accuracy			1.00	1876
macro avg	1.00	1.00	1.00	1876
weighted avg	1.00	1.00	1.00	1876

Accuracy: 0.9989339019189766

[231]: `import matplotlib.pyplot as plt`
`import seaborn as sns`
`from sklearn.metrics import confusion_matrix`


```

# Replace with bagging_model or rf_model as needed
y_pred = bagging_model.predict(X_test)

# Confusion matrix
cm = confusion_matrix(y_test, y_pred)
labels = bagging_model.classes_ # or rf_model.classes_

plt.figure(figsize=(8, 6))
sns.heatmap(cm, annot=True, fmt="d", cmap="Blues", xticklabels=labels,
            yticklabels=labels)
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title("Confusion Matrix - Bagging Classifier")
plt.show()

```



```

[ ]: # The confusion matrix above indicates that the model correctly classified all
      of the

```



```

# high risk patients, however, accidentally classified 2 of the moderate risk as
↳ low risk.
# Only misclassifying two patients is a fairly good model, and with health care
↳ data it is better to err on the side of more risk
# Which we can see that the model did

```

```

[232]: from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import confusion_matrix, classification_report,
↳ accuracy_score
import seaborn as sns
import matplotlib.pyplot as plt

# Train the Random Forest model
rf_model = RandomForestClassifier(n_estimators=100, random_state=42)
rf_model.fit(X_train, y_train)

# Predict on the test set
rf_preds = rf_model.predict(X_test)

# Confusion matrix
rf_cm = confusion_matrix(y_test, rf_preds)
print("Random Forest Classification Report:")
print(classification_report(y_test, rf_preds))
print("Accuracy:", accuracy_score(y_test, rf_preds))

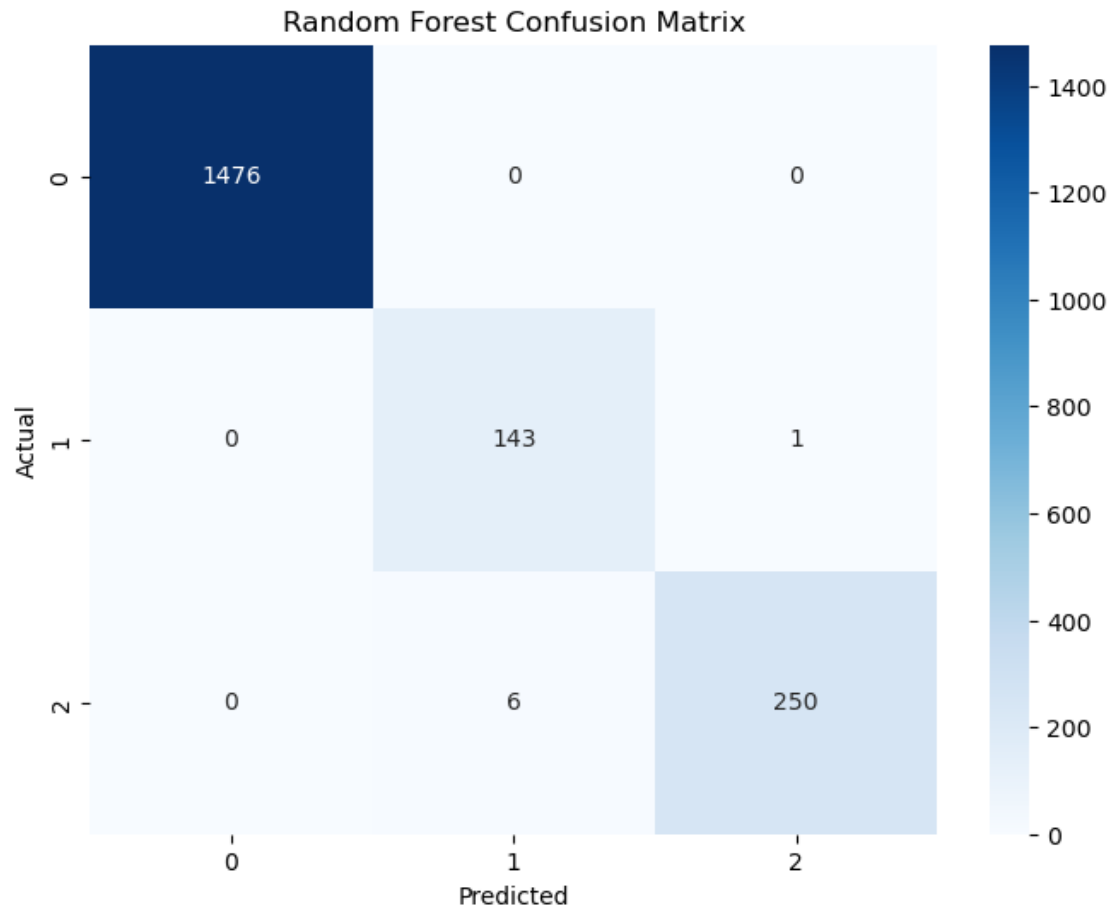
# Plot the confusion matrix
plt.figure(figsize=(8,6))
sns.heatmap(rf_cm, annot=True, fmt='d', cmap='Blues',
            xticklabels=rf_model.classes_, yticklabels=rf_model.classes_)
plt.title('Random Forest Confusion Matrix')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.show()

```

Random Forest Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	1476
1	0.96	0.99	0.98	144
2	1.00	0.98	0.99	256
accuracy			1.00	1876
macro avg	0.99	0.99	0.99	1876
weighted avg	1.00	1.00	1.00	1876

Accuracy: 0.996268656716418



```
[ ]: # The random forest classifier correctly labeled all of the high risk patients,␣  
      ↪however, it mixed up 6 of the low risk patients and categorized them in␣  
      ↪moderate risk.  
      # It also incorrectly labeled one patient as low risk when they should have␣  
      ↪been moderate.
```

```
[ ]: # Based off of the two models performance for classifying patients, I would␣  
      ↪recommend using the bagging classifier over the random forest method.
```

```
[ ]: # Github link: https://github.com/AjaStarkey/  
      ↪Masters-in-Data-Science-Class-Projects
```